



# Setting-up GNN Prediction Tasks

# GNN Training Pipeline (5)



(5) How do we split our dataset into train / validation / test set?



Dataset split

# Dataset Split: Fixed / Random Split

- **Fixed split:** We will split our dataset **once**
  - **Training set:** used for optimizing GNN parameters
  - **Validation set:** develop model/hyperparameters
  - **Test set:** held out until we report final performance
- **A concern:** sometimes we cannot guarantee that the test set will really be held out
- **Random split:** we will **randomly split** our dataset into training / validation / test
  - We report **average performance over different random seeds**

# Why Splitting Graphs is Special




- Suppose we want to split an image dataset
  - **Image classification:** Each data point is an image
  - Here **data points are independent**
    - Image 5 will not affect our prediction on image 1

Training

Validation

Test




# Why Splitting Graphs is Special



- **Splitting a graph dataset is different!**
  - **Node classification:** Each data point is a node
  - Here **data points are NOT independent**
    - **Node 5 will affect our prediction on node 1**, because it will participate in message passing → affect node 1's embedding

Training  
Validation  
Test



- **What are our options?**

# Why Splitting Graphs is Special




- **Solution 1 (Transductive setting):** The input graph can be observed in all the dataset splits (training, validation and test set).
- **We will only split the (node) labels**
  - At training time, we compute embeddings using the entire graph, and train using node 1&2's labels
  - At validation time, we compute embeddings using the entire graph, and evaluate on node 3&4's labels

Training

Validation

Test




# Why Splitting Graphs is Special

- **Solution 2 (Inductive setting):** We break the edges between splits **to get multiple graphs**
  - **Now we have 3 graphs that are independent.** Node 5 will not affect our prediction on node 1 any more
  - **At training time**, we compute embeddings **using the graph over node 1&2**, and train **using node 1&2's labels**
  - **At validation time**, we compute embeddings **using the graph over node 3&4**, and **evaluate on node 3&4's labels**

Training

Validation

Test



# Transductive / Inductive Settings


- **Transductive setting:** training / validation / test sets are **on the same graph**
  - The **dataset consists of one graph**
  - **The entire graph can be observed in all dataset splits, we only split the labels**
  - Only applicable to **node / edge** prediction tasks
- **Inductive setting:** training / validation / test sets are **on different graphs**
  - The **dataset consists of multiple graphs**
  - Each split can **only observe the graph(s) within the split**. A successful model should **generalize to unseen graphs**
  - Applicable to **node / edge / graph** tasks



# Example: Node Classification



- **Transductive node classification**
  - All the splits can observe the entire graph structure, but can only observe the labels of their respective nodes






Training

Validation

Test

- **Inductive node classification**
  - Suppose we have a dataset of 3 graphs
  - Each split contains an independent graph



Training

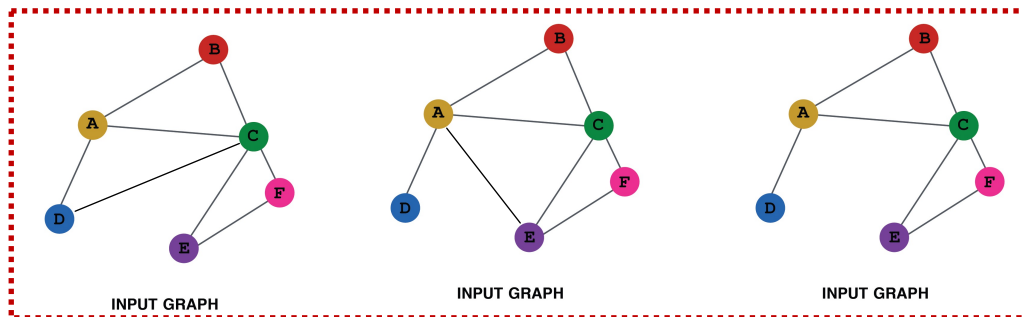
Validation

Test


# Example: Graph Classification




- Only the **inductive setting** is well defined for **graph classification**
  - Because **we have to test on unseen graphs**
  - Suppose we have a dataset of 5 graphs. Each split will contain independent graph(s).



**Training**



**Validation**




**Test**


# Example: Link Prediction




- **Goal of link prediction:** predict missing edges
- **Setting up link prediction is tricky:**
  - Link prediction is an unsupervised / self-supervised task. We need to **create the labels** and **dataset splits** on our own
  - Concretely, we need to **hide some edges from the GNN** and **let the GNN predict if the edges exist**



Original graph




Input graph to GNN



Predictions made by GNN

# Setting up Link Prediction




- **For link prediction, we will split edges twice**
- **Step 1: Assign 2 types of edges in the original graph**
  - **Message edges:** Used for GNN message passing
  - **Supervision edges:** Use for computing objectives
  - **After step 1:**
    - Only message edges will remain in the graph
    - Supervision edges are used as supervision for edge predictions made by the model, will not be fed into GNN!

# Setting up Link Prediction



- **Step 2: Split edges into train / validation / test**
- **Option 1: Inductive link prediction split**
  - Suppose we have a dataset of 3 graphs. Each inductive split will contain an independent graph






# Setting up Link Prediction



- **Step 2: Split edges into train / validation / test**
- **Option 1: Inductive link prediction split**
  - Suppose we have a dataset of 3 graphs. Each inductive split will contain an independent graph
  - In **train** or **val** or **test** set, each graph will have **2** types of edges: message edges + supervision edges
    - Supervision edges are not the input to GNN

Message  
edge ———


Supervision  
edge .....



# Setting up Link Prediction




- **Option 2: Transductive link prediction split:**
  - This is the default setting when people talk about link prediction
  - Suppose we have a dataset of 1 graph



# Setting up Link Prediction



- **Option 2: Transductive link prediction split:**
  - By definition of “transductive”, the entire graph can be observed in all dataset splits
    - But since edges are both part of graph structure and the supervision, we need to hold out **validation** / **test** edges
    - To train the **training** set, we further need to hold out **supervision edges** for the **training** set




- **Next:** we will show the exact settings




# Setting up Link Prediction




## ■ Option 2: Transductive link prediction split:




The original graph



(1) At training time:  
Use **training message edges** to predict **training supervision edges**



(2) At validation time:  
Use **training message edges & training supervision edges** to predict **validation edges**




(3) At test time:  
Use **training message edges & training supervision edges & validation edges** to predict **test edges**

# Setting up Link Prediction




## ■ Summary: Transductive link prediction split:



The original graph

Split  
→

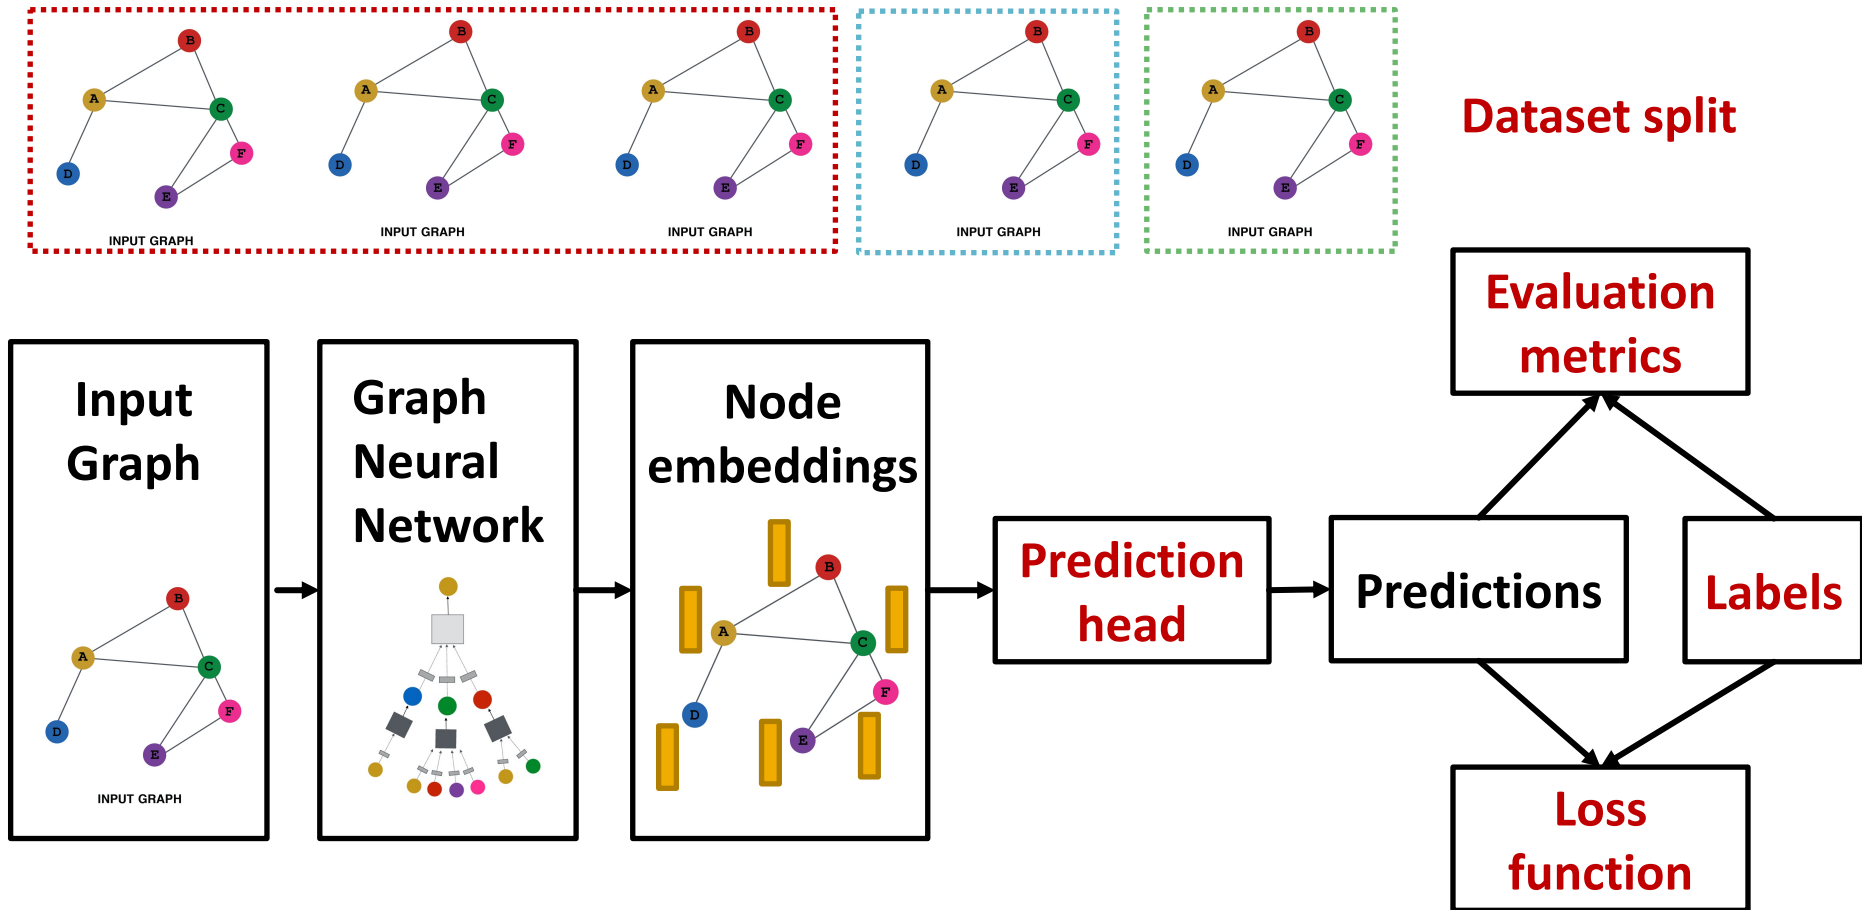


Split Graph with  
4 types of edges

Training message edges  
Training supervision edges  
Validation edges  
Test edges

- **Note:** Link prediction settings are tricky and complex. You may find papers do link prediction differently.
- Luckily, we have full support in **PyG** and [GraphGym](#)

# GNN Training Pipeline



## Implementation resources:

[DeepSNAP](#) provides core modules for this pipeline

[GraphGym](#) further implements the full pipeline to facilitate GNN design

# Summary of the Lecture



- **We introduce a general GNN framework:**
  - **GNN Layer:**
    - Transformation + Aggregation
    - Classic GNN layers: GCN, GraphSAGE, GAT
  - **Layer connectivity:**
    - The over-smoothing problem
    - Solution: skip connections
  - **Graph Augmentation:**
    - Feature augmentation
    - Structure augmentation
  - **Learning Objectives**
    - The full training pipeline of a GNN