



# **1. Basics of deep learning**



## **2. Deep learning for graphs**

## **3. Graph Convolutional Networks**

## **4. GNNs subsume CNNs and Transformers**

# Machine Learning as Optimization

- **Supervised learning:** we are given input  $x$ , and the goal is to predict label  $y$ .
- **Input  $x$  can be:**
  - Vectors of real numbers
  - Sequences (natural language)
  - Matrices (images)
  - Graphs (potentially with node and edge features)
- **We formulate the task as an optimization problem.**

# Machine Learning as Optimization

- **Formulate the task as an optimization problem:**

$$\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$$


**Objective function**

- $\Theta$ : a set of **parameters** we optimize
  - Could contain one or more scalars, vectors, matrices ...
  - E.g.  $\Theta = \{Z\}$  in the shallow encoder (the embedding lookup)

- $\mathcal{L}$ : **loss function**. Example: L2 loss

$$\mathcal{L}(\mathbf{y}, f(\mathbf{x})) = \|\mathbf{y} - f(\mathbf{x})\|_2$$

- Other common loss functions:
  - L1 loss, huber loss, max margin (hinge loss), cross entropy ...
  - See <https://pytorch.org/docs/stable/nn.html#loss-functions>

# Loss Function Example



- One common loss for classification: **cross entropy (CE)**
- Label  $\mathbf{y}$  is a categorical vector (**one-hot encoding**)
  - e.g.  $\mathbf{y} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \end{bmatrix}$  y is of class "3"

- $f(\mathbf{x}) = \text{Softmax}(g(\mathbf{x}))$

- Recall from lecture 3:  $f(\mathbf{x})_i = \frac{e^{g(\mathbf{x})_i}}{\sum_{j=1}^C e^{g(\mathbf{x})_j}}$

$g(\mathbf{x})_i$  denotes  $i$ -th coordinate of the vector output of func.  $g(\mathbf{x})$

where  $C$  is the number of classes.

- e.g.  $f(\mathbf{x}) = \begin{bmatrix} 0.1 & 0.3 & 0.4 & 0.1 & 0.1 \end{bmatrix}$

- $\text{CE}(\mathbf{y}, f(\mathbf{x})) = -\sum_{i=1}^C (y_i \log f(\mathbf{x})_i)$

- $y_i$  and  $f(\mathbf{x})_i$  are the **actual** and **predicted** values of the  $i$ -th class.
  - **Intuition:** the lower the loss, the closer the prediction is to one-hot

- **Total loss over all training examples:**

- $\mathcal{L} = \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{T}} \text{CE}(\mathbf{y}, f(\mathbf{x}))$

- $\mathcal{T}$ : training set containing all pairs of data and labels  $(\mathbf{x}, \mathbf{y})$


# Machine Learning as Optimization



- How to optimize the **objective function**?
- **Gradient vector**: Direction and rate of fastest increase

$$\nabla_{\Theta} \mathcal{L} = \left( \frac{\partial \mathcal{L}}{\partial \Theta_1}, \frac{\partial \mathcal{L}}{\partial \Theta_2}, \dots \right)$$

Partial derivative



<https://en.wikipedia.org/wiki/Gradient>

- $\Theta_1, \Theta_2 \dots$  : components of  $\Theta$
- Recall **directional derivative** of a multi-variable function (e.g.  $\mathcal{L}$ ) along a given vector represents the instantaneous rate of change of the function along the vector.
- **Gradient is the directional derivative in the direction of largest increase.**

# Gradient Descent



- **Iterative algorithm:** repeatedly update weights in the (opposite) direction of gradients until convergence

$$\Theta \leftarrow \Theta - \eta \nabla_{\Theta} \mathcal{L}$$

- **Training:** Optimize  $\Theta$  iteratively
  - **Iteration:** 1 step of gradient descent
- **Learning rate (LR)  $\eta$ :**
  - Hyperparameter that controls the size of gradient step
  - Can vary over the course of training (LR scheduling)
- **Ideal termination condition:** gradient = 0
  - In practice, we stop training if it no longer improves performance on **validation set** (part of dataset we hold out from training).

# Stochastic Gradient Descent (SGD)

## ■ Problem with gradient descent:

- Exact gradient requires computing  $\nabla_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$ , where  $\mathbf{x}$  is the **entire** dataset!
  - This means summing gradient contributions over all the points in the dataset
  - Modern datasets often contain billions of data points
  - Extremely expensive for every gradient descent step

## ■ Solution: Stochastic gradient descent (SGD)

- At every step, pick a different **minibatch**  $\mathcal{B}$  containing a subset of the dataset, use it as input  $\mathbf{x}$

# Minibatch SGD



## ■ Concepts:

- **Batch size**: the number of data points in a minibatch
  - E.g. number of nodes for node classification task
- **Iteration**: 1 step of SGD on a minibatch
- **Epoch**: one full pass over the dataset (# iterations is equal to ratio of dataset size and batch size)

## ■ SGD is unbiased estimator of full gradient:

- But there is no guarantee on the rate of convergence
- In practice often requires tuning of learning rate

## ■ Common optimizer that improves over SGD:

- Adam, Adagrad, Adadelata, RMSprop ...



# Neural Network Function



- **Objective:**  $\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$
- In deep learning, function  $f$  can be very complex
- **Example:**

- To start simple, consider linear function

$$f(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x}, \quad \Theta = \{\mathbf{W}\}$$

- Then, if  $f$  returns a scalar, then  $\mathbf{W}$  is a learnable **vector**

$$\nabla_{\mathbf{W}} f = \left( \frac{\partial f}{\partial w_1}, \frac{\partial f}{\partial w_2}, \frac{\partial f}{\partial w_3} \dots \right)$$

- But, if  $f$  returns a vector, then  $\mathbf{W}$  is the **weight matrix**

$$\nabla_{\mathbf{W}} f = \begin{bmatrix} \frac{\partial f_1}{\partial w_{11}} & \frac{\partial f_1}{\partial w_{12}} \\ \frac{\partial f_2}{\partial w_{21}} & \frac{\partial f_2}{\partial w_{22}} \end{bmatrix}$$

Jacobian  
matrix of  $f$

# Intuition: Back Propagation



- **Goal:**  $\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$

- To minimize  $\mathcal{L}$ , we need to evaluate the gradient:

$$\nabla_{\mathbf{w}} \mathcal{L} = \left( \frac{\partial \mathcal{L}}{\partial w_1}, \frac{\partial \mathcal{L}}{\partial w_2}, \frac{\partial \mathcal{L}}{\partial w_3} \dots \right)$$

which means we need to derive derivative of  $\mathcal{L}$ .

- **Overview of Back-propagation:**

- $\mathcal{L}$  is composed from some set of predefined building block functions  $g(\cdot)$
- For each such  $g$  we also have its derivative  $g'$
- Then we can automatically compute  $\nabla_{\Theta} \mathcal{L}$  by evaluating appropriate funcs.  $g'$  on the minibatch  $\mathcal{B}$ .

# Back-propagation



- How about a more complex function:

$$f(\mathbf{x}) = W_2(W_1\mathbf{x}), \Theta = \{W_1, W_2\}$$

In other words:

$$f(\mathbf{x}) = W_2(W_1\mathbf{x})$$

$$h(\mathbf{x}) = W_1\mathbf{x}$$

$$g(z) = W_2z$$

- Recall **chain rule**:

$$\frac{df}{dx} = \frac{dg}{dh} \cdot \frac{dh}{dx} \text{ or } f'(x) = g'(h(x))h'(x)$$

- **Example:**  $\nabla_{\mathbf{x}} f = \frac{\partial f}{\partial (W_1\mathbf{x})} \cdot \frac{\partial (W_1\mathbf{x})}{\partial \mathbf{x}}$


- **Back-propagation:** Use of **chain rule** to propagate gradients of intermediate steps, and finally obtain gradient of  $\mathcal{L}$  w.r.t.  $\Theta$ .

# Back-propagation Example (1)



- **Example:** Simple 2-layer linear network

- $f(\mathbf{x}) = g(h(\mathbf{x})) = W_2(W_1\mathbf{x})$



- $\mathcal{L} = \sum_{(\mathbf{x}, \mathbf{y}) \in \mathcal{B}} \left\| (\mathbf{y}, -f(\mathbf{x})) \right\|_2$

- The loss  $\mathcal{L}$  sums the L2 loss in a minibatch  $\mathcal{B}$ .

- **Hidden layer:**

- Intermediate representation of input  $\mathbf{x}$
  - Here we use  $h(\mathbf{x}) = W_1\mathbf{x}$  to denote the hidden layer
  - $f(\mathbf{x}) = W_2h(\mathbf{x})$

# Back-propagation Example (2)



## ■ Forward propagation:

Compute loss starting from input




Remember:

$$f(x) = W_2(W_1x)$$

$$h(x) = W_1x$$

$$g(z) = W_2z$$



## ■ Back-propagation to compute gradient of

$$\Theta = \{W_1, W_2\}$$

Start from loss, compute the gradient

$$\frac{\partial \mathcal{L}}{\partial W_2} = \frac{\partial \mathcal{L}}{\partial f} \cdot \frac{\partial f}{\partial W_2}$$

→  
Compute backwards

$$\frac{\partial \mathcal{L}}{\partial W_1} = \frac{\partial \mathcal{L}}{\partial f} \cdot \frac{\partial f}{\partial W_2} \cdot \frac{\partial W_2}{\partial W_1}$$

→  
Compute backwards

# Non-linearity




- Note that in  $f(\mathbf{x}) = W_2(W_1\mathbf{x})$ ,  $W_2W_1$  is another matrix (vector, if we do binary classification)
  - Hence  $f(\mathbf{x})$  is still linear w.r.t.  $\mathbf{x}$  no matter how many weight matrices we compose

- **We introduce non-linearity:**


- **Rectified linear unit (ReLU)**

$$\text{ReLU}(x) = \max(x, 0)$$



- **Sigmoid**

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




# Multi-layer Perceptron (MLP)



- Each layer of MLP combines linear transformation and non-linearity:

$$\mathbf{x}^{(l+1)} = \sigma(W_l \mathbf{x}^{(l)} + b^l)$$

- where  $W_l$  is weight matrix that transforms hidden representation at layer  $l$  to layer  $l + 1$
- $b^l$  is bias at layer  $l$ , and is added to the linear transformation of  $\mathbf{x}$
- $\sigma$  is non-linearity function (e.g., sigmoid)
- Suppose  $\mathbf{x}$  is 2-dimensional, with entries  $x_1$  and  $x_2$



Every layer:  
Linear transformation +  
non-linearity

# Summary



- **Objective function:**

$$\min_{\Theta} \mathcal{L}(\mathbf{y}, f(\mathbf{x}))$$

- $f$  can be a simple linear layer, an MLP, or other neural networks (e.g., a GNN later)
- Sample a minibatch of input  $\mathbf{x}$
- **Forward propagation:** Compute  $\mathcal{L}$  given  $\mathbf{x}$
- **Back-propagation:** Obtain gradient  $\nabla_{\mathbf{w}} \mathcal{L}$  using a chain rule.
- Use **stochastic gradient descent (SGD)** to optimize for  $\Theta$  over many iterations.